

dl-lab-exp-no-2

May 9, 2025

```
[1]: import numpy as np
import tensorflow as tf
from tensorflow.keras.datasets import imdb
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Embedding, LSTM, Dropout
from tensorflow.keras.optimizers import Adam
import matplotlib.pyplot as plt
```

```
[2]: (X_train, y_train), (X_test, y_test) = imdb.load_data(num_words=10000)
max_length = 200
X_train = pad_sequences(X_train, maxlen=max_length)
X_test = pad_sequences(X_test, maxlen=max_length)
print(f"X_train shape: {X_train.shape}")
print(f"X_test shape: {X_test.shape}")
```

```
X_train shape: (25000, 200)
X_test shape: (25000, 200)
```

```
[3]: model = Sequential()
model.add(Embedding(input_dim=10000, output_dim=128))
model.add(LSTM(units=128, dropout=0.2, recurrent_dropout=0.2))
model.add(Dropout(0.5))
model.add(Dense(1, activation='sigmoid'))
model.compile(optimizer='adam', loss='binary_crossentropy',
              metrics=['accuracy'])
model.summary()
```

```
Model: "sequential"
```

Layer (type)	Output Shape	
Param #		
embedding (Embedding)	?	0
(unbuilt)		

```

lstm (LSTM)                ?                0
↳(unbuilt)

dropout (Dropout)          ?                0
↳ 0

dense (Dense)              ?                0
↳(unbuilt)

```

Total params: 0 (0.00 B)

Trainable params: 0 (0.00 B)

Non-trainable params: 0 (0.00 B)

```

[4]: history = model.fit(X_train, y_train, epochs=3, batch_size=64,
↳validation_data=(X_test, y_test))

```

```

Epoch 1/3
391/391          247s 614ms/step -
accuracy: 0.6883 - loss: 0.5713 - val_accuracy: 0.8099 - val_loss: 0.4219
Epoch 2/3
391/391          207s 530ms/step -
accuracy: 0.8571 - loss: 0.3483 - val_accuracy: 0.8542 - val_loss: 0.3476
Epoch 3/3
391/391          191s 487ms/step -
accuracy: 0.8914 - loss: 0.2800 - val_accuracy: 0.8419 - val_loss: 0.3787


```


```

[5]: plt.plot(history.history['accuracy'], label='accuracy')
plt.plot(history.history['val_accuracy'], label = 'val_accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.ylim([0, 1])
plt.legend(loc='lower right')
plt.title('Model Accuracy')
plt.show()

plt.plot(history.history['loss'], label='loss')
plt.plot(history.history['val_loss'], label = 'val_loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend(loc='upper right')
plt.title('Model Loss')
plt.show()

```





```
[6]: test_loss, test_acc = model.evaluate(X_test, y_test, verbose=2)
      print(f"Test accuracy: {test_acc}")
```

```
782/782 - 40s - 51ms/step - accuracy: 0.8419 - loss: 0.3787
Test accuracy: 0.841920018196106
```

```
[7]: from sklearn.metrics import classification_report, confusion_matrix
      import seaborn as sns
```

```
[8]: # Get predictions
      y_pred = (model.predict(X_test) >= 0.5).astype(int).flatten()
      y_test = y_test.flatten()
```

```
782/782          40s 50ms/step
```

```
[9]: # Classification report
      print("\nClassification Report:")
      print(classification_report(y_test, y_pred, target_names=['Negative',
      ↪ 'Positive']))
```

Classification Report:

	precision	recall	f1-score	support
Negative	0.84	0.84	0.84	12500
Positive	0.84	0.84	0.84	12500
accuracy			0.84	25000
macro avg	0.84	0.84	0.84	25000
weighted avg	0.84	0.84	0.84	25000

```
[10]: # Confusion matrix
cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(6, 4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=['Negative', 'Positive'], yticklabels=['Negative', 'Positive'])
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.show()
```

