

DL Lab Exp No.3

```
[1]: import tensorflow as tf
from sklearn.model_selection import train_test_split
from mlxtend.plotting import plot_confusion_matrix
from sklearn import metrics
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
from tqdm.notebook import tqdm
import random
import warnings
warnings.filterwarnings("ignore")

[2]: (trainX, trainY), (testX, testY) = tf.keras.datasets.fashion_mnist.load_data()

trainX = trainX.reshape((trainX.shape[0], 28, 28, 1))
testX = testX.reshape((testX.shape[0], 28, 28, 1))

trainY_cat = tf.keras.utils.to_categorical(trainY)
testY_cat = tf.keras.utils.to_categorical(testY)

[3]: train_norm = trainX.astype('float32')
test_norm = testX.astype('float32')

train_norm = train_norm / 255.0
test_norm = test_norm / 255.0

[4]: class_names = ['T-shirt/top', 'Trouser', 'Pullover', 'Dress', 'Coat',
                    'Sandal', 'Shirt', 'Sneaker', 'Bag', 'Ankle boot']

[5]: plt.figure(figsize=(10,10))
for i in range(25):
    plt.subplot(5,5,i+1)
    plt.xticks([])
    plt.yticks([])
    plt.grid(False)
```

```
plt.imshow(trainX[i], cmap=plt.cm.binary)
plt.xlabel(class_names[trainY[i]])
plt.show()
```



```
[6]: model = tf.keras.models.Sequential([
    tf.keras.layers.Conv2D(64, kernel_size=(3, 3), input_shape=(28, 28, 1),
    ↪activation='relu', padding='same', name='conv-layer-1'),
    tf.keras.layers.AvgPool2D(pool_size=(2, 2), name='pooling-layer-1'),
    tf.keras.layers.Conv2D(32, kernel_size=(3, 3), activation='relu',
    ↪padding='same', name='conv-layer-2'),
    tf.keras.layers.AvgPool2D(pool_size=(2, 2), name='pooling-layer-2'),
    tf.keras.layers.GlobalAveragePooling2D(name='pooling-layer-3'),
```

```
tf.keras.layers.Dense(len(class_names), activation="softmax",  
↳name="output-layer")  
])
```

```
[7]: model.compile(loss="categorical_crossentropy", optimizer="adam",  
↳metrics=["accuracy"])
```

```
[8]: history = model.fit(trainX, trainY_cat, epochs=10, validation_data=(testX,  
↳testY_cat))
```

```
Epoch 1/10  
1875/1875          80s 41ms/step -  
accuracy: 0.6166 - loss: 1.2695 - val_accuracy: 0.7968 - val_loss: 0.5972  
Epoch 2/10  
1875/1875          48s 25ms/step -  
accuracy: 0.8084 - loss: 0.5460 - val_accuracy: 0.8129 - val_loss: 0.5320  
Epoch 3/10  
1875/1875          46s 25ms/step -  
accuracy: 0.8274 - loss: 0.4839 - val_accuracy: 0.8433 - val_loss: 0.4574  
Epoch 4/10  
1875/1875          47s 25ms/step -  
accuracy: 0.8415 - loss: 0.4441 - val_accuracy: 0.8323 - val_loss: 0.4671  
Epoch 5/10  
1875/1875          47s 25ms/step -  
accuracy: 0.8473 - loss: 0.4259 - val_accuracy: 0.8508 - val_loss: 0.4301  
Epoch 6/10  
1875/1875          56s 30ms/step -  
accuracy: 0.8624 - loss: 0.3952 - val_accuracy: 0.8601 - val_loss: 0.3975  
Epoch 7/10  
1875/1875          49s 26ms/step -  
accuracy: 0.8679 - loss: 0.3760 - val_accuracy: 0.8607 - val_loss: 0.3909  
Epoch 8/10  
1875/1875          51s 27ms/step -  
accuracy: 0.8735 - loss: 0.3594 - val_accuracy: 0.8621 - val_loss: 0.3979  
Epoch 9/10  
1875/1875          53s 28ms/step -  
accuracy: 0.8766 - loss: 0.3504 - val_accuracy: 0.8757 - val_loss: 0.3546  
Epoch 10/10  
1875/1875          45s 24ms/step -  
accuracy: 0.8834 - loss: 0.3331 - val_accuracy: 0.8786 - val_loss: 0.3491
```

```
[9]: tf.keras.utils.plot_model(model, show_shapes=True)
```

You must install graphviz (see instructions at
<https://graphviz.gitlab.io/download/>) for `plot_model` to work.

```
[10]: model.summary()
```

Model: "sequential"

Layer (type) ↳Param #	Output Shape	
conv-layer-1 (Conv2D) ↳640	(None, 28, 28, 64)	↳
pooling-layer-1 (AveragePooling2D) ↳ 0	(None, 14, 14, 64)	↳
conv-layer-2 (Conv2D) ↳18,464	(None, 14, 14, 32)	↳
pooling-layer-2 (AveragePooling2D) ↳ 0	(None, 7, 7, 32)	↳
pooling-layer-3 ↳ 0 (GlobalAveragePooling2D) ↳	(None, 32)	↳
output-layer (Dense) ↳330	(None, 10)	↳

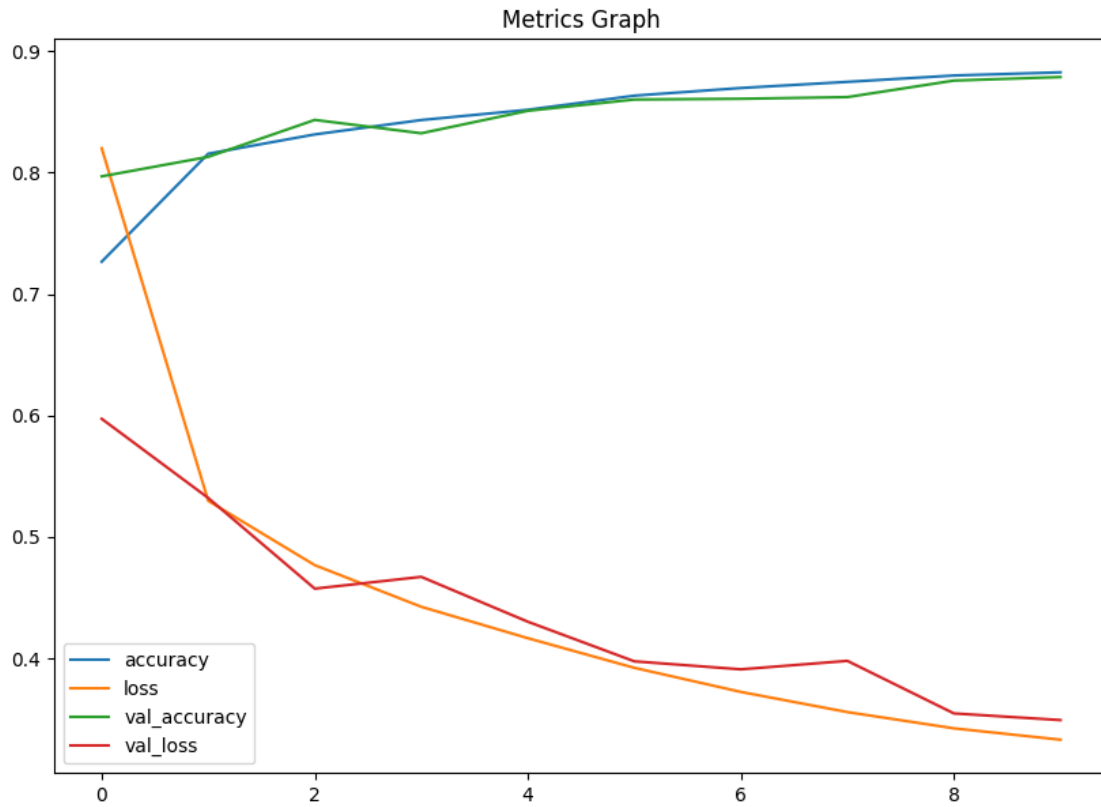
Total params: 58,304 (227.75 KB)

Trainable params: 19,434 (75.91 KB)

Non-trainable params: 0 (0.00 B)

Optimizer params: 38,870 (151.84 KB)

```
[11]: pd.DataFrame(history.history).plot(figsize=(10,7))
plt.title("Metrics Graph")
plt.show()
```



```
[12]: model.evaluate(testX, testY_cat)
```

```
313/313          3s 9ms/step -
accuracy: 0.8810 - loss: 0.3418
```

```
[12]: [0.3491450846195221, 0.878600001335144]
```

```
[13]: predictions = model.predict(testX)
```

```
313/313          3s 7ms/step
```

```
[14]: predictions = tf.argmax(predictions, axis=1)
```

```
[15]: tf.keras.utils.plot_model(model, show_shapes=True)
```

You must install graphviz (see instructions at <https://graphviz.gitlab.io/download/>) for `plot_model` to work.

```
[16]: model.evaluate(testX, testY_cat)
```

```
313/313          3s 10ms/step -
accuracy: 0.8810 - loss: 0.3418
```

```
[16]: [0.3491450846195221, 0.878600001335144]
```

```
[17]: predictions = model.predict(testX)
      predictions = tf.argmax(predictions, axis=1)
      y_test = tf.argmax(testY_cat, axis=1)
      y_test = tf.Variable(y_test)
```

313/313 2s 8ms/step


```
[18]: print("Accuracy: ", metrics.accuracy_score(y_test, predictions))
```

Accuracy: 0.8786

```
[19]: print(metrics.classification_report(y_test, predictions))
```

	precision	recall	f1-score	support
0	0.88	0.76	0.81	1000
1	0.97	0.97	0.97	1000
2	0.83	0.81	0.82	1000
3	0.85	0.87	0.86	1000
4	0.72	0.88	0.79	1000
5	0.97	0.97	0.97	1000
6	0.69	0.65	0.67	1000
7	0.93	0.98	0.95	1000
8	0.98	0.97	0.97	1000
9	0.98	0.94	0.96	1000
accuracy			0.88	10000
macro avg	0.88	0.88	0.88	10000
weighted avg	0.88	0.88	0.88	10000

```
[20]: cm = metrics.confusion_matrix(y_test, predictions)
      plot_confusion_matrix(cm, figsize=(10,7), class_names=class_names)
      plt.title("Confusion Matrix")
      plt.show()
```



```
[21]: images = []
labels = []
random_indices = random.sample(range(len(testX)), 10)
for idx in random_indices:
    images.append(testX[idx])
    labels.append(testY_cat[idx])
images = np.array(images)
labels = np.array(labels)

fig = plt.figure(figsize=(20, 8))
rows = 2
cols = 5
```

```

x = 1
for image, label in zip(images, labels):
    fig.add_subplot(rows, cols, x)
    prediction = model.predict(tf.expand_dims(image, axis=0))
    prediction = class_names[tf.argmax(prediction.flatten())]
    label = class_names[tf.argmax(label)]
    plt.title(f"Label: {label}, Prediction: {prediction}")
    plt.imshow(image/255.)
    plt.axis("off")
    x += 1

```

```

1/1          0s 108ms/step
1/1          0s 101ms/step
1/1          0s 99ms/step
1/1          0s 160ms/step
1/1          0s 124ms/step
1/1          0s 136ms/step
1/1          0s 101ms/step
1/1          0s 106ms/step
1/1          0s 94ms/step
1/1          0s 93ms/step

```

